

Assignment2(1828251)

May 27, 2022

1 ACML Assignment 2

2 Ziyaad Ballim (1828251)

2.1 Importing Libraries

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from prettytable import PrettyTable
import seaborn as sns
sns.set()
```

3 Reading in the data

This dataset has been collected from patients tested for diabetes. There are 8 features namely: -Pregnancies

-Glucose

-BloodPressure

-SkinThickness

-Insulin

-BMI

-DiabetesPedigreeFunction

-Age

and one output column “Ouput” which represents the presence of diabetes or not with 1 or 0 respectively.

```
[ ]: data = pd.read_csv('diabetes-dataset.csv')
```

```
[ ]: data.describe()
```

```
[ ]:      Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin  \
count  2000.000000  2000.000000  2000.000000  2000.000000  2000.000000
mean    3.703500    121.182500    69.145500    20.935000    80.254000
std     3.306063     32.068636    19.188315    16.103243    111.180534
min     0.000000     0.000000     0.000000     0.000000     0.000000
25%     1.000000     99.000000    63.500000     0.000000     0.000000
50%     3.000000    117.000000    72.000000    23.000000    40.000000
75%     6.000000    141.000000    80.000000    32.000000    130.000000
max     17.000000    199.000000    122.000000    110.000000    744.000000
```

```
      BMI  DiabetesPedigreeFunction      Age      Outcome
count  2000.000000      2000.000000  2000.000000  2000.000000
mean    32.193000      0.470930    33.090500    0.342000
std     8.149901      0.323553    11.786423    0.474498
min     0.000000      0.078000    21.000000    0.000000
25%    27.375000      0.244000    24.000000    0.000000
50%    32.300000      0.376000    29.000000    0.000000
75%    36.800000      0.624000    40.000000    1.000000
max    80.600000      2.420000    81.000000    1.000000
```

```
[ ]: data.head()
```

```
[ ]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0         2         138         62           35         0  33.6
1         0         84         82           31        125  38.2
2         0        145          0           0         0  44.2
3         0        135         68           42        250  42.3
4         1        139         62           41        480  40.7

      DiabetesPedigreeFunction  Age  Outcome
0         0.127         47         1
1         0.233         23         0
2         0.630         31         1
3         0.365         24         1
4         0.536         21         0
```

Checking for null values

```
[ ]: data.isna().sum()
```

```
[ ]: Pregnancies      0
      Glucose         0
      BloodPressure    0
      SkinThickness     0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
```

```
Age                                0
Outcome                            0
dtype: int64
```

3.1 Functions for splitting and standardizing the data

Splitting data into training, validation and testing subsets

```
[ ]: def SplitData(x,y,testSize):
    trainx,testx, trainy,testy =
    ↪train_test_split(x,y,test_size=testSize,random_state=40)
    return trainx,trainy,testx,testy
```

Standardizing the data

```
[ ]: def StandardizeData(data):
    rows,cols = np.shape(data)

    for i in range(cols):
        Feature = data[:,i]
        mean = np.mean(Feature)
        standard_deviation = np.std(Feature)
        Feature -= mean
        Feature /= standard_deviation
        data[:,i] = Feature

    return data
```

Scaling

```
[ ]: def ScaleData(data):
    scaledx = StandardizeData(data)
    return scaledx
```

```
[ ]: def Split_Data(data,validation_and_test_size,scale):
    arrData = data.to_numpy()
    np.random.shuffle(arrData)

    rows,cols = np.shape(arrData)
    Features = arrData[:, :cols-1]
    Target = arrData[:, -1]
    Target = np.reshape(Target, (np.shape(Target)[0],1))
    if(scale):
        Features = ScaleData(Features)
    test_size = np.round(validation_and_test_size*0.3,2)
    trainx,trainy,tempx,tempy
    ↪=SplitData(Features,Target,validation_and_test_size)
    validationx,validationy,testx,testy =SplitData(tempx,tempy,test_size)
```

```
return trainx,validationx,testx, trainy,validationy,testy
```

3.2 Splitting and standardizing the data

```
[ ]: trainx,validationx,testx,trainy,validationy,testy = Split_Data(data,0.3,True)
```

```
[ ]: print(len(trainx))
      print(len(validationx))
      print(len(testx))
      print(len(trainy))
      print(len(validationy))
      print(len(testy))
```

```
1400
546
54
1400
546
54
```

3.3 The functions below are used to construct, train and test the intended Neural Networks

Initialising the weights of the neural network between 0 and 1

```
[ ]: def InitWeights(NumFeatures,NumHidLayers,NumNeurons,NumOutputNeurons):
      InLayerWeights = np.random.rand(NumNeurons[0],NumFeatures+1)
      Weights = [InLayerWeights]
      for i in range(NumHidLayers):
          W = None
          if(NumHidLayers -1 == i):
              W = np.random.rand(NumOutputNeurons,NumNeurons[i]+1)

          else:
              W = np.random.rand(NumNeurons[i+1],NumNeurons[i]+1)

          Weights.append(W)

      return Weights
```

Converting the output probabilities to output classes

```
[ ]: def ProbToClass(LatestActivation):
      Indices = np.where(LatestActivation<0.5)
      LatestActivation[Indices] = 0
      Indices2 = np.where(LatestActivation>=0.5)
      LatestActivation[Indices2] = 1
      return LatestActivation
```

Controlling the use of the different activation functions

```
[ ]: def ActFunction(Z,Activation):
    a = None
    deriv_z = None
    if (Activation == "sigmoid"):
        a = (1/(1+np.exp(-Z)))
        deriv_z = a * (1-a)
    elif (Activation == "tanh"):
        a = np.tanh(Z)
        deriv_z = 1 - (a**2)
    elif (Activation == "relu"):
        a = np.maximum(0,Z)
        deriv_z = np.copy(a)
        deriv_z[deriv_z<=0] = 0
        deriv_z[deriv_z>0] = 1
    return a,deriv_z
```

Forward propogation in order to obtain predicted outputs

```
[ ]: def ForwardProp(features,weights,activations):
    A_Vals = [features]
    Derivs = [0]
    n = len(weights)
    for i in range(n):
        z = None
        aprev = np.insert(A_Vals[i],0,1,axis=1)
        z = np.dot(weights[i],aprev.T)
        A_Vals[i] = aprev
        a,derivative_of_z = ActFunction(z,activations[i])
        A_Vals.append(a.T)
        Derivs.append(derivative_of_z)

    return A_Vals,Derivs
```

Computing the errors

```
[ ]: def Errors(a,y,Weights,derivs):
    n = len(a)
    errors = [a[n-1]-y]
    loopLength = n -2
    for j in range(loopLength):
        e = (((Weights[loopLength - j].T)[1:,:])@(errors[j].T))*
        ↪derivs[loopLength - j]
        errors.append(e.T)

    return errors
```

Obtaining gradients

```
[ ]: def Grads(Errors,Activations):  
    grads = []  
    reversed_errors = list(reversed(Errors))  
    n = len(Activations)  
    for i in range(n-1):  
        g = (reversed_errors[i].T)@(Activations[i])  
        grads.append(g)  
    return grads
```

Performing back propogations

```
[ ]: def   
    ↪BackProp(Features,Target,Weights,Activations,numDataPoints,regularization,gradients):  
    ↪  
        avalues,derivatives = ForwardProp(Features,Weights,Activations)  
  
        errors = Errors(avales,Target,Weights,derivatives)  
  
        Updatedgradients = Grads(errors,avalues)  
        for k in range(len(gradients)):  
            gradients[k]+=Updatedgradients[k]  
  
        if(regularization != 0):  
            n = len(gradients)  
            for i in range(n):  
                gradients[i] = (1/numDataPoints)*gradients[i] +   
    ↪regularization*Weights[i]  
        else:  
            n = len(gradients)  
            for i in range(n):  
                gradients[i] = (1/numDataPoints)*gradients[i]  
  
    return gradients
```

Checking for the convergence of the weights

```
[ ]: def WeightConv(NewWeights,OldWeights):  
    n = len(NewWeights)  
    Controls = []  
    for k in range(n):  
        if(np.linalg.norm(NewWeights[k]-OldWeights[k])<0.00005):  
            Controls.append(True)  
        else:  
            Controls.append(False)  
  
    numFalse = np.where(np.asarray(Controls)==False)
```

```

length = len(numFalse[0])
if(length>1):
    return False
else:
    return True

```

Dealing with division errors

```

[ ]: def DivError(x):
    ETA = 0.0000000001
    return np.maximum(x, ETA)

```

Calculating the costs

```

[ ]: def Cost(Target,predicted,regularization,weights):
    n = len(Target)
    yhat_inv = np.subtract(1.0,predicted)
    y_inv = np.subtract(1.0,Target)
    yhat = DivError(predicted)
    yhat_inv = DivError(yhat_inv)
    weightsSummation=0;
    for w in weights:
        weightsSummation+= np.sum(w**2)
    loss = -1/n * (np.sum(np.multiply(np.log(yhat), Target) + np.
↪multiply((y_inv), np.log(yhat_inv))))+ (regularization/
↪(2*n))*weightsSummation
    return loss

```

Performing gradient descent

```

[ ]: def ↵
↪GradDescent(Features,Target,weights,numDataPoints,learningRate,regularization,Epochs,Activa
↪
    gradients = [None]*len(weights)
    loss = []
    for i in range(len(gradients)):
        gradients[i] = np.zeros(np.shape(weights[i]))

    iterations=0
    control = True
    while iterations<Epochs and control:
        Old_Weights = weights.copy()
        D = ↵
↪BackProp(Features,Target,weights,ActivationFunctions,numDataPoints,regularization,gradients

        for j in range(len(D)):

```

```

        weights[j] = weights[j]- learningRate*D[j]

    if(WeightConv(weights,Old_Weights)):
        control = False
    iterations+=1
    a,v = ForwardProp(Features,weights,ActivationFunctions)
    predictedValues = a[len(a)-1]
    lossValue = Cost(Target,predictedValues,regularization,weights)
    loss.append(lossValue)
    return weights,iterations,loss

```

Finding the output predictions

```

[ ]: def Predict(testx,testy,weights,ActivationFunctions):
    act,der = ForwardProp(testx,weights,ActivationFunctions)
    act[len(act)-1] = ProbToClass(act[len(act)-1])
    predictedValues = act[len(act)-1]
    conf = confusion_matrix(testy,predictedValues)
    acc = accuracy_score(testy,predictedValues)
    print(f'The accuracy is: {acc*100}%')
    print("Confusion Matrix: ")
    sns.heatmap(conf,annot=True,fmt='g')
    plt.show()
    return acc

```

Plotting the error graph over time

```

[ ]: def ErrorGraph(error):
    x = np.arange(1,(len(error)+1))
    sns.lineplot(x = x ,y = error)
    plt.title("Error vs Epochs")
    plt.xlabel("Epochs")
    plt.ylabel("Cost")
    plt.show()

```

Fitting the neural network using the above functions and getting the necessary input

```

[ ]: def FitNN(TrainingX,TrainingY,NumberOfHiddenLayers,NumberOfNeuronsPerHiddenLayer,NumberOfNeuronsPerOutputLayer):
    ↪
    ↪
    Number_of_datapoints, Number_of_features = np.shape(TrainingX)

    Weight_Parameters = []
    ↪InitWeights(Number_of_features,NumberOfHiddenLayers,NumberOfNeuronsPerHiddenLayer,NumberOfNeuronsPerOutputLayer)
    assert(len(ActivationFunctions) ==
    ↪NumberOfHiddenLayers+NumberOfNeuronsInOutputLayer), "Please ensure that the
    ↪number of activation functions specified is the same as the layers!"

```



```

Weights,epochs_to_converge,error = ␣
↪GradDescent(TrainingX,TrainingY,Weight_Parameters,Number_of_datapoints,learningRate,regular
ErrorGraph(error)
weights = Weights
activationFunctions = ActivationFunctions
return Weights,epochs_to_converge,error, activationFunctions

```

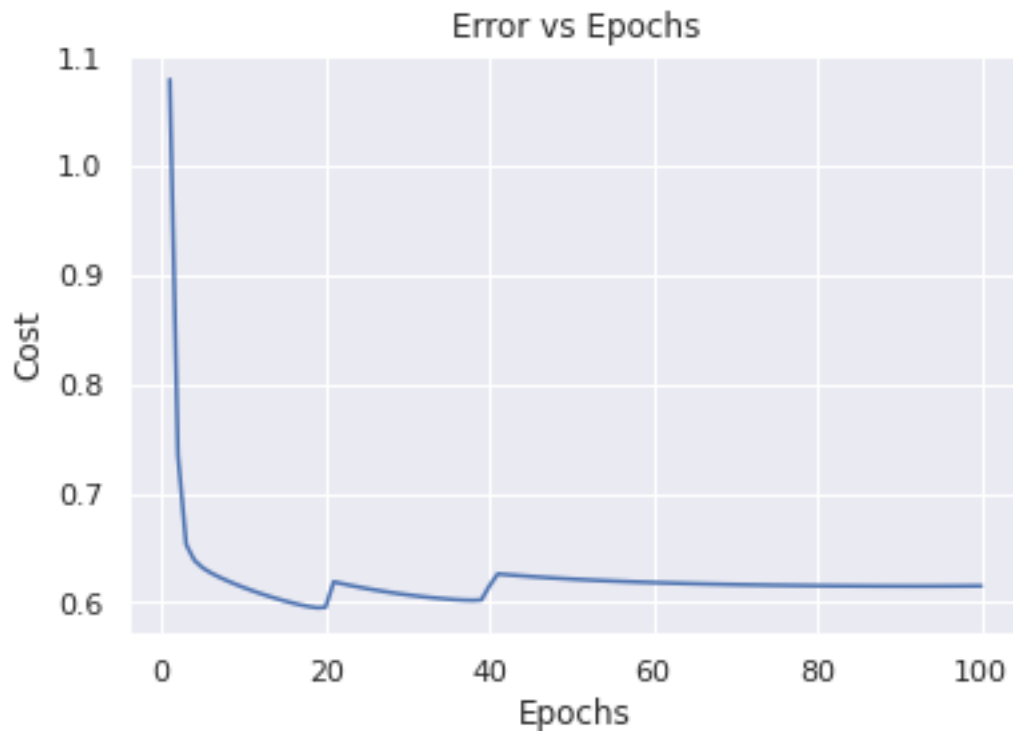
3.4 Implementation of Neural network and training the network using back-propagation

3.4.1 Relu Activation Function

```

[ ]: ReluWeights, ReluEpochs, ReluError, ReluActivations = ␣
↪FitNN(trainx,trainy,2,[2,2],1,["relu","relu","relu"],0.1,0,100)
print("Number of epochs until convergence:"+str(ReluEpochs)+"\n")
print("Training Accuracy")
Relu_Acc=Predict(trainx,trainy,ReluWeights,ReluActivations)
print("Validation Accuracy")
Relu_Acc_Val=Predict(validationx,validationy,ReluWeights,ReluActivations)
print("Test Accuracy")
Relu_Acc_Test=Predict(testx,testy,ReluWeights,ReluActivations)

```




Number of epochs until convergence:100

Training Accuracy

The accuracy is: 72.64285714285714%


Confusion Matrix:



Validation Accuracy

The accuracy is: 74.54212454212454%


Confusion Matrix:



Test Accuracy


The accuracy is: 70.37037037037037%

Confusion Matrix:



3.4.2 Sigmoid Activation Function

```
[ ]: SigWeights, SigEpochs, SigError, SigActivations =   
    ↳FitNN(trainx,trainy,2,[2,2],1,["sigmoid","sigmoid","sigmoid"],0.1,0,100)  
print("Number of epochs until convergence:"+str(SigEpochs)+"\n")  
print("Training Accuracy")  
Sig_Acc=Predict(trainx,trainy,SigWeights,SigActivations)  
print("Validation Accuracy")  
Sig_Acc_Val=Predict(validationx,validationy,SigWeights,SigActivations)  
print("Test Accuracy")  
Sig_Acc_Test=Predict(testx,testy,SigWeights,SigActivations)
```




Number of epochs until convergence:100


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:



Test Accuracy


The accuracy is: 51.85185185185185%

Confusion Matrix:



3.4.3 Tanh Activation Function

```
[ ]: TanhWeights, TanhEpochs, TanhError, TanhActivations =   
    ↪ FitNN(trainx, trainy, 2, [2, 2], 1, ["tanh", "tanh", "tanh"], 0.1, 0, 100)  
    print("Number of epochs until convergence:" + str(TanhEpochs) + "\n")  
    print("Training Accuracy")  
    Tanh_Acc = Predict(trainx, trainy, TanhWeights, TanhActivations)  
    print("Validation Accuracy")  
    Tanh_Acc_Val = Predict(validationx, validationy, TanhWeights, TanhActivations)  
    print("Test Accuracy")  
    Tanh_Acc_Test = Predict(testx, testy, TanhWeights, TanhActivations)
```




Number of epochs until convergence:100


Training Accuracy

The accuracy is: 69.35714285714286%

Confusion Matrix:




Validation Accuracy
The accuracy is: 71.97802197802197%
Confusion Matrix:



Test Accuracy

The accuracy is: 68.51851851851852%

Confusion Matrix:



```
[ ]: InitialResults = PrettyTable(["Activation Function","Epochs to train","Training Accuracy",  
    "Validation Accuracy","Testing Accuracy"])
```

```
InitialResults.add_row(["ReLU",ReluEpochs,np.round(Relu_Acc*100,2),np.  
    round(Relu_Acc_Val*100,2),np.round(Relu_Acc_Test*100,2)])  
InitialResults.add_row(["Sigmoid",SigEpochs,np.round(Sig_Acc*100,2),np.  
    round(Sig_Acc_Val*100,2),np.round(Sig_Acc_Test*100,2)])  
InitialResults.add_row(["Tanh",TanhEpochs,np.round(Tanh_Acc*100,2),np.  
    round(Tanh_Acc_Val*100,2),np.round(Tanh_Acc_Test*100,2)])  
print(InitialResults)
```

```
+-----+-----+-----+-----+  
+-----+  
| Activation Function | Epochs to train | Training Accuracy | Validation  
Accuracy | Testing Accuracy |  
+-----+-----+-----+-----+  
+-----+  
|      ReLu      | |      100      | |      72.64     | |      74.54     |  
|      70.37     | |               | |               | |               |  
|      Sigmoid   | |      100      | |      65.29     | |      68.5      |  
|      51.85     | |               | |               | |               |  
|      Tanh      | |      100      | |      69.36     | |      71.98     |
```

```
|      68.52      |
+-----+-----+-----+-----+
--+-----+-----+-----+-----+
```

3.5 Testing different network sizes(Number of hidden layers and number of neurons)

3.5.1 Relu

3.5.2 1 Hidden Layers vs 2 Hidden Layers

8 Neurons per layer vs 4 Neurons per layer

```
[ ]: print("1 Layer with 8 Neurons")
R18Weights, R18Epochs, R18Error, R18Activations = \
    ↪FitNN(trainx,trainy,1,[8],1,["relu","relu"],0.1,0,1000)
print("Number of epochs until convergence:"+str(R18Epochs)+"\n")
print("Training Accuracy")
R18_Acc=Predict(trainx,trainy,R18Weights,R18Activations)
print("Validation Accuracy")
R18_Acc_Val=Predict(validationx,validationy,R18Weights,R18Activations)
print("Test Accuracy")
R18_Acc_Test=Predict(testx,testy,R18Weights,R18Activations)

print("1 Layer with 4 Neurons")
R110Weights, R110Epochs, R110Error, R110Activations = \
    ↪FitNN(trainx,trainy,1,[4],1,["relu","relu"],0.1,0,1000)
print("Number of epochs until convergence:"+str(R110Epochs)+"\n")
print("Training Accuracy")
R110_Acc=Predict(trainx,trainy,R110Weights,R110Activations)
print("Validation Accuracy")
R110_Acc_Val=Predict(validationx,validationy,R110Weights,R110Activations)
print("Test Accuracy")
R110_Acc_Test=Predict(testx,testy,R110Weights,R110Activations)

print("2 Layers with 8 Neurons")
R28Weights, R28Epochs, R28Error, R28Activations = \
    ↪FitNN(trainx,trainy,2,[8,8],1,["relu","relu","relu"],0.1,0,1000)
print("Number of epochs until convergence:"+str(R28Epochs)+"\n")
print("Training Accuracy")
R28_Acc=Predict(trainx,trainy,R28Weights,R28Activations)
print("Validation Accuracy")
R28_Acc_Val=Predict(validationx,validationy,R28Weights,R28Activations)
print("Test Accuracy")
R28_Acc_Test=Predict(testx,testy,R28Weights,R28Activations)

print("2 Layers with 4 Neurons")
R210Weights, R210Epochs, R210Error, R210Activations = \
    ↪FitNN(trainx,trainy,2,[4,4],1,["relu","relu","relu"],0.1,0,1000)
```

```


print("Number of epochs until convergence:"+str(R210Epochs)+"\n")
print("Training Accuracy")
R210_Acc=Predict(trainx,trainy,R210Weights,R210Activations)
print("Validation Accuracy")
R210_Acc_Val=Predict(validationx,validationy,R210Weights,R210Activations)
print("Test Accuracy")
R210_Acc_Test=Predict(testx,testy,R210Weights,R210Activations)

ReluGen = PrettyTable(["Hidden Layers","Neurons in layers","Epochs to_
↳train","Training Accuracy","Validation Accuracy","Testing Accuracy"])

ReluGen.add_row(["1","8",R18Epochs,np.round(R18_Acc*100,2),np.
↳round(R18_Acc_Val*100,2),np.round(R18_Acc_Test*100,2)])
ReluGen.add_row(["1","4",R110Epochs,np.round(R110_Acc*100,2),np.
↳round(R110_Acc_Val*100,2),np.round(R110_Acc_Test*100,2)])
ReluGen.add_row(["2","8",R28Epochs,np.round(R28_Acc*100,2),np.
↳round(R28_Acc_Val*100,2),np.round(R28_Acc_Test*100,2)])
ReluGen.add_row(["2","4",R210Epochs,np.round(R210_Acc*100,2),np.
↳round(R210_Acc_Val*100,2),np.round(R210_Acc_Test*100,2)])
print(ReluGen)

```

1 Layer with 8 Neurons




Number of epochs until convergence:1000

Training Accuracy

The accuracy is: 76.0%


Confusion Matrix:



Validation Accuracy

The accuracy is: 76.55677655677655%


Confusion Matrix:




Test Accuracy

The accuracy is: 74.07407407407408%

Confusion Matrix:



1 Layer with 4 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 76.07142857142857%

Confusion Matrix:




Validation Accuracy
The accuracy is: 76.92307692307693%
Confusion Matrix:




Test Accuracy

The accuracy is: 74.07407407407408%

Confusion Matrix:



2 Layers with 8 Neurons




Number of epochs until convergence:4


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



2 Layers with 4 Neurons




Number of epochs until convergence:42


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:



Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



| Hidden Layers | | Neurons in layers | | Epochs to train | | Training Accuracy | | Validation Accuracy | | Testing Accuracy | |
|---------------|--|-------------------|--|-----------------|--|-------------------|--|---------------------|--|------------------|--|
| 1 | | 8 | | 1000 | | 76.0 | | | | | |
| 76.56 | | 74.07 | | | | | | | | | |
| 1 | | 4 | | 1000 | | 76.07 | | | | | |
| 76.92 | | 74.07 | | | | | | | | | |
| 2 | | 8 | | 4 | | 65.29 | | | | | |
| 68.5 | | 51.85 | | | | | | | | | |
| 2 | | 4 | | 42 | | 65.29 | | | | | |
| 68.5 | | 51.85 | | | | | | | | | |

3.5.3 Sigmoid

1 layer vs 2 layers

8 Neurons vs 9 Neurons

```

[ ]: print("1 Layer with 8 Neurons")
S18Weights, S18Epochs, S18Error, S18Activations =
    ↳FitNN(trainx,trainy,1,[8],1,["sigmoid","sigmoid"],0.1,0,1000)
print("Number of epochs until convergence:"+str(S18Epochs)+"\n")
print("Training Accuracy")
S18_Acc=Predict(trainx,trainy,S18Weights,S18Activations)
print("Validation Accuracy")
S18_Acc_Val=Predict(validationx,validationy,S18Weights,S18Activations)
print("Test Accuracy")
S18_Acc_Test=Predict(testx,testy,S18Weights,S18Activations)

print("1 Layer with 9 Neurons")
S110Weights, S110Epochs, S110Error, S110Activations =
    ↳FitNN(trainx,trainy,1,[9],1,["sigmoid","sigmoid"],0.1,0,1000)
print("Number of epochs until convergence:"+str(S110Epochs)+"\n")
print("Training Accuracy")
S110_Acc=Predict(trainx,trainy,S110Weights,S110Activations)
print("Validation Accuracy")
S110_Acc_Val=Predict(validationx,validationy,S110Weights,S110Activations)
print("Test Accuracy")
S110_Acc_Test=Predict(testx,testy,S110Weights,S110Activations)

print("2 Layers with 8 Neurons")
S28Weights, S28Epochs, S28Error, S28Activations =
    ↳FitNN(trainx,trainy,2,[8,8],1,["sigmoid","sigmoid","sigmoid"],0.1,0,1000)
print("Number of epochs until convergence:"+str(S28Epochs)+"\n")
print("Training Accuracy")
S28_Acc=Predict(trainx,trainy,S28Weights,S28Activations)
print("Validation Accuracy")
S28_Acc_Val=Predict(validationx,validationy,S28Weights,S28Activations)
print("Test Accuracy")
S28_Acc_Test=Predict(testx,testy,S28Weights,S28Activations)

print("2 Layers with 9 Neurons")
S210Weights, S210Epochs, S210Error, S210Activations =
    ↳FitNN(trainx,trainy,2,[9,9],1,["sigmoid","sigmoid","sigmoid"],0.1,0,1000)
print("Number of epochs until convergence:"+str(S210Epochs)+"\n")
print("Training Accuracy")
S210_Acc=Predict(trainx,trainy,S210Weights,S210Activations)
print("Validation Accuracy")
S210_Acc_Val=Predict(validationx,validationy,S210Weights,S210Activations)
print("Test Accuracy")
S210_Acc_Test=Predict(testx,testy,S210Weights,S210Activations)

SigmoidGen = PrettyTable(["Hidden Layers","Neurons in layers","Epochs to
    ↳train","Training Accuracy","Validation Accuracy","Testing Accuracy"])


```

```

SigmoidGen.add_row(["1", "8", S18Epochs, np.round(S18_Acc*100,2), np.
    ↳round(S18_Acc_Val*100,2), np.round(S18_Acc_Test*100,2)])
SigmoidGen.add_row(["1", "9", S110Epochs, np.round(S110_Acc*100,2), np.
    ↳round(S110_Acc_Val*100,2), np.round(S110_Acc_Test*100,2)])
SigmoidGen.add_row(["2", "8", S28Epochs, np.round(S28_Acc*100,2), np.
    ↳round(S28_Acc_Val*100,2), np.round(S28_Acc_Test*100,2)])
SigmoidGen.add_row(["2", "9", S210Epochs, np.round(S210_Acc*100,2), np.
    ↳round(S210_Acc_Val*100,2), np.round(S210_Acc_Test*100,2)])
print(SigmoidGen)

```

1 Layer with 8 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 76.57142857142857%

Confusion Matrix:




Validation Accuracy
The accuracy is: 77.83882783882784%
Confusion Matrix:




Test Accuracy

The accuracy is: 83.33333333333334%

Confusion Matrix:



1 Layer with 9 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 76.42857142857142%

Confusion Matrix:




Validation Accuracy
The accuracy is: 79.3040293040293%
Confusion Matrix:




Test Accuracy

The accuracy is: 83.33333333333334%

Confusion Matrix:



2 Layers with 8 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



2 Layers with 9 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 72.57142857142857%

Confusion Matrix:




Validation Accuracy
The accuracy is: 74.9084249084249%
Confusion Matrix:



Test Accuracy

The accuracy is: 72.22222222222221%

Confusion Matrix:



| Hidden Layers | Neurons in layers | Epochs to train | Training Accuracy | Validation Accuracy | Testing Accuracy |
|---------------|-------------------|-----------------|-------------------|---------------------|------------------|
| | | | | | |
| 1 | 8 | 1000 | 76.57 | | |
| 77.84 | 83.33 | | | | |
| 1 | 9 | 1000 | 76.43 | | |
| 79.3 | 83.33 | | | | |
| 2 | 8 | 1000 | 65.29 | | |
| 68.5 | 51.85 | | | | |
| 2 | 9 | 1000 | 72.57 | | |
| 74.91 | 72.22 | | | | |

3.5.4 Tanh

1 layer vs 2 layers

3 Neurons per layer vs 18 Neurons per layer

```

[ ]: print("1 Layer with 3 Neurons")
T18Weights, T18Epochs, T18Error, T18Activations =
    ↳FitNN(trainx,trainy,1,[3],1,["tanh","tanh"],0.1,0,1000)
print("Number of epochs until convergence:"+str(T18Epochs)+"\n")
print("Training Accuracy")
T18_Acc=Predict(trainx,trainy,T18Weights,T18Activations)
print("Validation Accuracy")
T18_Acc_Val=Predict(validationx,validationy,T18Weights,T18Activations)
print("Test Accuracy")
T18_Acc_Test=Predict(testx,testy,T18Weights,T18Activations)

print("1 Layer with 18 Neurons")
T110Weights, T110Epochs, T110Error, T110Activations =
    ↳FitNN(trainx,trainy,1,[18],1,["tanh","tanh"],0.1,0,1000)
print("Number of epochs until convergence:"+str(T110Epochs)+"\n")
print("Training Accuracy")
T110_Acc=Predict(trainx,trainy,T110Weights,T110Activations)
print("Validation Accuracy")
T110_Acc_Val=Predict(validationx,validationy,T110Weights,T110Activations)
print("Test Accuracy")
T110_Acc_Test=Predict(testx,testy,T110Weights,T110Activations)

print("2 Layers with 3 Neurons")
T28Weights, T28Epochs, T28Error, T28Activations =
    ↳FitNN(trainx,trainy,2,[3,3],1,["tanh","tanh","tanh"],0.1,0,1000)
print("Number of epochs until convergence:"+str(T28Epochs)+"\n")
print("Training Accuracy")
T28_Acc=Predict(trainx,trainy,T28Weights,T28Activations)
print("Validation Accuracy")
T28_Acc_Val=Predict(validationx,validationy,T28Weights,T28Activations)
print("Test Accuracy")
T28_Acc_Test=Predict(testx,testy,T28Weights,T28Activations)

print("2 Layers with 18 Neurons")
T210Weights, T210Epochs, T210Error, T210Activations =
    ↳FitNN(trainx,trainy,2,[18,18],1,["tanh","tanh","tanh"],0.1,0,1000)
print("Number of epochs until convergence:"+str(T210Epochs)+"\n")
print("Training Accuracy")
T210_Acc=Predict(trainx,trainy,T210Weights,T210Activations)
print("Validation Accuracy")
T210_Acc_Val=Predict(validationx,validationy,T210Weights,T210Activations)
print("Test Accuracy")
T210_Acc_Test=Predict(testx,testy,T210Weights,T210Activations)

TanhGen = PrettyTable(["Hidden Layers","Neurons in layers","Epochs to
    ↳train","Training Accuracy","Validation Accuracy","Testing Accuracy"])


```

```

TanhGen.add_row(["1", "3", T18Epochs, np.round(T18_Acc*100,2), np.
    ↳round(T18_Acc_Val*100,2), np.round(T18_Acc_Test*100,2)])
TanhGen.add_row(["1", "18", T110Epochs, np.round(T110_Acc*100,2), np.
    ↳round(T110_Acc_Val*100,2), np.round(T110_Acc_Test*100,2)])
TanhGen.add_row(["2", "3", T28Epochs, np.round(T28_Acc*100,2), np.
    ↳round(T28_Acc_Val*100,2), np.round(T28_Acc_Test*100,2)])
TanhGen.add_row(["2", "18", T210Epochs, np.round(T210_Acc*100,2), np.
    ↳round(T210_Acc_Val*100,2), np.round(T210_Acc_Test*100,2)])
print(TanhGen)

```

1 Layer with 3 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 77.28571428571429%

Confusion Matrix:




Validation Accuracy
The accuracy is: 78.3882783882784%
Confusion Matrix:




Test Accuracy

The accuracy is: 85.18518518518519%

Confusion Matrix:



1 Layer with 18 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 78.57142857142857%

Confusion Matrix:




Validation Accuracy
The accuracy is: 78.93772893772893%
Confusion Matrix:




Test Accuracy

The accuracy is: 87.03703703703704%

Confusion Matrix:



2 Layers with 3 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 76.28571428571429%

Confusion Matrix:




Validation Accuracy
The accuracy is: 76.73992673992674%
Confusion Matrix:




Test Accuracy

The accuracy is: 81.48148148148148%

Confusion Matrix:



2 Layers with 18 Neurons




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 70.07142857142857%


Confusion Matrix:



Validation Accuracy
The accuracy is: 69.78021978021978%
Confusion Matrix:



Confusion Matrix:



| +-----+-----+-----+ | | | | |
|---------------------|--|-------------------|--|-----------------|
| +-----+-----+ | | | | |
| Hidden Layers | | Neurons in layers | | Epochs to train |
| Validation Accuracy | | Testing Accuracy | | |
| +-----+-----+-----+ | | | | |
| +-----+-----+ | | | | |
| 1 | | 3 | | 1000 |
| 78.39 | | 85.19 | | 77.29 |
| 1 | | 18 | | 1000 |
| 78.94 | | 87.04 | | 78.57 |
| 2 | | 3 | | 1000 |
| 76.74 | | 81.48 | | 76.29 |
| 2 | | 18 | | 1000 |
| 69.78 | | 66.67 | | 70.07 |
| +-----+-----+-----+ | | | | |
| +-----+-----+ | | | | |

4 Discussing results

```
[ ]: print("Initial results with 2 layers with 2 neurons each")
      print(InitialResults)
```

Initial results with 2 layers with 2 neurons each

| Initial results with 2 layers with 2 neurons each | | | | |
|--|---------|--|-----|--|
| ----- | | | | |
| --+ | | | | |
| Activation Function Epochs to train Training Accuracy Validation | | | | |
| Accuracy Testing Accuracy | | | | |
| +-----+ | | | | |
| --+ | | | | |
| | ReLU | | 100 | |
| | 70.37 | | | |
| | Sigmoid | | 100 | |
| | 51.85 | | | |
| | Tanh | | 100 | |
| | 68.52 | | | |
| +-----+ | | | | |
| --+ | | | | |

```
[ ]: print("Relu activation function with different network sizes")
      print(ReluGen)
```

Relu activation function with different network sizes

| Relu activation function with different network sizes | | | | |
|---|---|--|-------|--|
| ----- | | | | |
| -----+ | | | | |
| Hidden Layers Neurons in layers Epochs to train Training Accuracy | | | | |
| Validation Accuracy Testing Accuracy | | | | |
| +-----+ | | | | |
| -----+ | | | | |
| | 1 | | 8 | |
| 76.56 | | | 74.07 | |
| | 1 | | 4 | |
| 76.92 | | | 74.07 | |
| | 2 | | 8 | |
| 68.5 | | | 51.85 | |
| | 2 | | 4 | |
| 68.5 | | | 51.85 | |
| +-----+ | | | | |
| -----+ | | | | |

```
[ ]: print("Sigmoid activation function with different network sizes")
      print(SigmoidGen)
```

Sigmoid activation function with different network sizes

| Sigmoid activation function with different network sizes | | | | |
|--|--|--|--|--|
| ----- | | | | |
| -----+ | | | | |

| Hidden Layers | Neurons in layers | Epochs to train | Training Accuracy | Validation Accuracy | Testing Accuracy |
|---------------|-------------------|-----------------|-------------------|---------------------|------------------|
| 1 | 8 | 1000 | 76.57 | 77.84 | 83.33 |
| 1 | 9 | 1000 | 76.43 | 79.3 | 83.33 |
| 2 | 8 | 1000 | 65.29 | 68.5 | 51.85 |
| 2 | 9 | 1000 | 72.57 | 74.91 | 72.22 |

```
[ ]: print("Tanh activation function with different network sizes")
      print(TanhGen)
```

Tanh activation function with different network sizes

| Hidden Layers | Neurons in layers | Epochs to train | Training Accuracy | Validation Accuracy | Testing Accuracy |
|---------------|-------------------|-----------------|-------------------|---------------------|------------------|
| 1 | 3 | 1000 | 77.29 | 78.39 | 85.19 |
| 1 | 18 | 1000 | 78.57 | 78.94 | 87.04 |
| 2 | 3 | 1000 | 76.29 | 76.74 | 81.48 |
| 2 | 18 | 1000 | 70.07 | 69.78 | 66.67 |

There has been a use of three activation functions namely: -Relu -Sigmoid -Tanh These have been used to examine the different effects that they have on the same data. When looking at the initial accuracies when using 2 hidden layers with 2 nodes each, the tanh activation function seems to out perform the other two for the training and validation, however the Relu and Sigmoid functions do the best when it comes to the hidden test data. Therefore implying that the relu and sigmoid functions are less fitted to the training data.

Thereafter, for each activation function we have looked at the effect of different network sizes. Each activation function has been provided with different network sizes that seem to provide optimal results.

The Relu function with 1 hidden layer consisting of 8 nodes performed the best overall when averaging the accuracies over the training, validation and testing data.

The learning rate and regularization has been consistent throughout, so we shall now study the effect of these using the best performing network from the previous observations.


5 Different learning rates

```
[ ]: LearningResults = PrettyTable(["Learning Rate","Epochs to train","Training_
    ↳Accuracy","Validation Accuracy","Testing Accuracy"])

for i in range(10):
    print("Learning Rate:"+str(1/10**i))
    R18Weights, R18Epochs, R18Error, R18Activations =_
    ↳FitNN(trainx,trainy,1,[8],1,["relu","relu"],1/10**i,0,1000)
    print("Number of epochs until convergence:"+str(R18Epochs)+"\n")
    print("Training Accuracy")
    R18_Acc=Predict(trainx,trainy,R18Weights,R18Activations)
    print("Validation Accuracy")
    R18_Acc_Val=Predict(validationx,validationy,R18Weights,R18Activations)
    print("Test Accuracy")
    R18_Acc_Test=Predict(testx,testy,R18Weights,R18Activations)
    LearningResults.add_row([1/10**i,R18Epochs,np.round(R18_Acc*100,2),np.
    ↳round(R18_Acc_Val*100,2),np.round(R18_Acc_Test*100,2)])

print(LearningResults)
```

Learning Rate:1.0




Number of epochs until convergence:18


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Learning Rate:0.1




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 74.92857142857143%

Confusion Matrix:




Validation Accuracy
The accuracy is: 76.007326007326%
Confusion Matrix:




Test Accuracy

The accuracy is: 74.07407407407408%

Confusion Matrix:



Learning Rate:0.01




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 74.21428571428571%

Confusion Matrix:




Validation Accuracy
The accuracy is: 75.64102564102564%
Confusion Matrix:




Test Accuracy

The accuracy is: 77.7777777777779%

Confusion Matrix:



Learning Rate:0.001




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 67.42857142857143%

Confusion Matrix:




Validation Accuracy
 The accuracy is: 67.76556776556777%
 Confusion Matrix:




Test Accuracy

The accuracy is: 59.25925925925925%

Confusion Matrix:



Learning Rate:0.0001




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 34.785714285714285%

Confusion Matrix:




Validation Accuracy
 The accuracy is: 34.065934065934066%
 Confusion Matrix:




Test Accuracy

The accuracy is: 48.148148148145%

Confusion Matrix:



Learning Rate:1e-05




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 57.85714285714286%

Confusion Matrix:




Validation Accuracy
The accuracy is: 55.311355311355314%
Confusion Matrix:




Test Accuracy

The accuracy is: 59.25925925925925%

Confusion Matrix:



Learning Rate:1e-06




Number of epochs until convergence:1


Training Accuracy

The accuracy is: 34.714285714285715%

Confusion Matrix:




Validation Accuracy
The accuracy is: 31.5018315018315%
Confusion Matrix:



Test Accuracy

The accuracy is: 48.148148148145%

Confusion Matrix:



Learning Rate:1e-07




Number of epochs until convergence:1


Training Accuracy

The accuracy is: 34.714285714285715%

Confusion Matrix:




Validation Accuracy
The accuracy is: 31.5018315018315%
Confusion Matrix:



Test Accuracy

The accuracy is: 48.148148148145%

Confusion Matrix:



Learning Rate:1e-08




Number of epochs until convergence:1


Training Accuracy

The accuracy is: 34.714285714285715%

Confusion Matrix:




Validation Accuracy
The accuracy is: 31.5018315018315%
Confusion Matrix:




Test Accuracy

The accuracy is: 48.148148148148145%

Confusion Matrix:



Learning Rate:1e-09




Number of epochs until convergence:1


Training Accuracy

The accuracy is: 57.92857142857143%

Confusion Matrix:




Validation Accuracy
The accuracy is: 56.41025641025641%
Confusion Matrix:



Test Accuracy

The accuracy is: 57.407407407405%

Confusion Matrix:



| Learning Rate | | Epochs to train | Training Accuracy | | Validation Accuracy |
|------------------|------|-----------------|-------------------|-------|---------------------|
| Testing Accuracy | | | | | |
| | | | | | |
| 1.0 | 18 | 65.29 | 68.5 | 51.85 | |
| 0.1 | 1000 | 74.93 | 76.01 | 74.07 | |
| 0.01 | 1000 | 74.21 | 75.64 | 77.78 | |
| 0.001 | 1000 | 67.43 | 67.77 | 59.26 | |
| 0.0001 | 1000 | 34.79 | 34.07 | 48.15 | |
| 1e-05 | 1000 | 57.86 | 55.31 | 59.26 | |
| 1e-06 | 1 | 34.71 | 31.5 | 48.15 | |
| 1e-07 | 1 | 34.71 | 31.5 | 48.15 | |

| | | | | | | | | |
|---------------------------------|-------|--|---|--|-------|--|-------|--|
| | 1e-08 | | 1 | | 34.71 | | 31.5 | |
| 48.15 | | | | | | | | |
| | 1e-09 | | 1 | | 57.93 | | 56.41 | |
| 57.41 | | | | | | | | |
| +-----+-----+-----+-----+-----+ | | | | | | | | |
| -----+ | | | | | | | | |

As seen by the table above, the learning rates of 0.1 and 0.01 produce the best results. Once the learning rate is 0.000001 or smaller, the results converge after 1 epoch, due to the fact that the learning rate is too low.


6 Different Regularization

```
[ ]: RegResults = PrettyTable(["Regularization","Epochs to train","Training_
    ↳Accuracy","Validation Accuracy","Testing Accuracy"])

for i in range(20):
    print("Regularization:"+str(0.1*i))
    R18Weights, R18Epochs, R18Error, R18Activations =_
    ↳FitNN(trainx,trainy,1,[8],1,["relu","relu"],0.1,0.1*i,1000)
    print("Number of epochs until convergence:"+str(R18Epochs)+"\n")
    print("Training Accuracy")
    R18_Acc=Predict(trainx,trainy,R18Weights,R18Activations)
    print("Validation Accuracy")
    R18_Acc_Val=Predict(validationx,validationy,R18Weights,R18Activations)
    print("Test Accuracy")
    R18_Acc_Test=Predict(testx,testy,R18Weights,R18Activations)
    RegResults.add_row([0.1*i,R18Epochs,np.round(R18_Acc*100,2),np.
    ↳round(R18_Acc_Val*100,2),np.round(R18_Acc_Test*100,2)])

print(RegResults)
```

Regularization:0.0




Number of epochs until convergence:1000


Training Accuracy

The accuracy is: 77.35714285714286%

Confusion Matrix:




Validation Accuracy
The accuracy is: 78.2051282051282%
Confusion Matrix:




Test Accuracy

The accuracy is: 77.7777777777779%

Confusion Matrix:



Regularization:0.1




Number of epochs until convergence:499


Training Accuracy

The accuracy is: 74.71428571428571%


Confusion Matrix:




Validation Accuracy
The accuracy is: 75.27472527472527%
Confusion Matrix:



Test Accuracy
The accuracy is: 72.2222222222221%
Confusion Matrix:



Regularization:0.2




Number of epochs until convergence:404


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:0.30000000000000004




Number of epochs until convergence:771


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:0.4




Number of epochs until convergence:342


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:0.5




Number of epochs until convergence:155


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:0.6000000000000001




Number of epochs until convergence:129


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:0.7000000000000001




Number of epochs until convergence:164


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:0.8




Number of epochs until convergence:135


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:0.9




Number of epochs until convergence:114


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.0




Number of epochs until convergence:98


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.1




Number of epochs until convergence:72


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.2000000000000002




Number of epochs until convergence:66


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.3




Number of epochs until convergence:63


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.4000000000000001




Number of epochs until convergence:52


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.5




Number of epochs until convergence:49


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.6




Number of epochs until convergence:52


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.7000000000000002




Number of epochs until convergence:44


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.8




Number of epochs until convergence:47


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:




Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



Regularization:1.9000000000000001




Number of epochs until convergence:42


Training Accuracy

The accuracy is: 65.28571428571428%

Confusion Matrix:




Validation Accuracy
The accuracy is: 68.4981684981685%
Confusion Matrix:



Test Accuracy

The accuracy is: 51.85185185185185%

Confusion Matrix:



| -----+ | | | | |
|---------------------|------------------|-----------------|-------------------|------------|
| --+- | | | | |
| Regularization | | Epochs to train | Training Accuracy | Validation |
| Accuracy | Testing Accuracy | | | |
| -----+ | | | | |
| --+- | | | | |
| 0.0 | | 1000 | 77.36 | 78.21 |
| 77.78 | | | | |
| 0.1 | | 499 | 74.71 | 75.27 |
| 72.22 | | | | |
| 0.2 | | 404 | 65.29 | 68.5 |
| 51.85 | | | | |
| 0.30000000000000004 | | 771 | 65.29 | 68.5 |
| 51.85 | | | | |
| 0.4 | | 342 | 65.29 | 68.5 |
| 51.85 | | | | |
| 0.5 | | 155 | 65.29 | 68.5 |
| 51.85 | | | | |
| 0.6000000000000001 | | 129 | 65.29 | 68.5 |
| 51.85 | | | | |
| 0.7000000000000001 | | 164 | 65.29 | 68.5 |
| 51.85 | | | | |

| | | | | | | | |
|---------------------------|--------------------|--|-----|--|-------|--|------|
| | 0.8 | | 135 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 0.9 | | 114 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.0 | | 98 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.1 | | 72 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.2000000000000002 | | 66 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.3 | | 63 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.4000000000000001 | | 52 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.5 | | 49 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.6 | | 52 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.7000000000000002 | | 44 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.8 | | 47 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| | 1.9000000000000001 | | 42 | | 65.29 | | 68.5 |
| | 51.85 | | | | | | |
| +-----+-----+-----+-----+ | | | | | | | |
| --+-----+ | | | | | | | |

The implementation of regularization seems to have had a negative effect on the accuracy of the network. When the regularization is set to 0, the network performs the best.